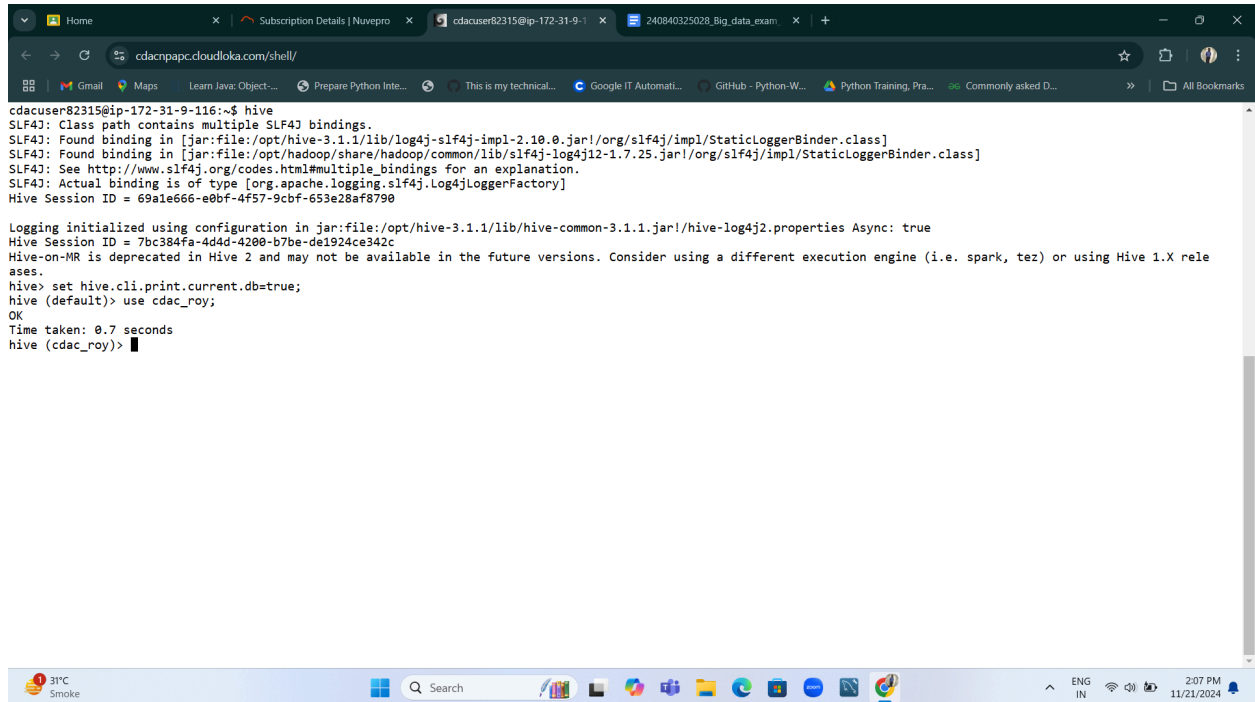


Big_data_exam: 240840325028

Q1.) 1.)



The screenshot shows a terminal window with a dark background. The terminal displays the output of the 'hive' command. It starts with 'cdacuser82315@ip-172-31-9-116:~\$ hive'. The output includes SLF4J warnings about multiple bindings, logging initialization details, and the execution of 'set hive.cli.print.current.db=true;' and 'use cdac_roy;'. The terminal also shows the system tray at the bottom with a temperature of 31°C, a search bar, and various application icons.

```
cdacuser82315@ip-172-31-9-116:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/opt/hive-3.1.1/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/opt/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 69a1e666-e0bf-4f57-9cbf-653e28af8798

Logging initialized using configuration in jar:file:/opt/hive-3.1.1/lib/hive-common-3.1.1.jar!/hive-log4j2.properties Async: true
Hive Session ID = 7bc384fa-4d4d-4200-b7be-de1924ce342c
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> set hive.cli.print.current.db=true;
hive (default)> use cdac_roy;
OK
Time taken: 0.7 seconds
hive (cdac_roy)>
```

```
select a1.name,a2.name from airport a1 join routes r on
a1.airport_id = r.src_airport_id join airport a2 on
a2.airport_id = r.dest_airport_id limit 10;
```

```
cdacnpac.cloudloka.com/shell/

Number of reduce tasks not specified. Defaulting to jobconf value of: 4
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1732089968849_2334, Tracking URL = http://master:6318/proxy/applica
tion_1732089968849_2334/
Kill Command = /opt/hadoop/bin/mapred job -kill job_1732089968849_2334
Hadoop job information for Stage-2: number of mappers: 3; number of reducers: 4
2024-11-21 08:56:03,288 Stage-2 map = 0%, reduce = 0%
2024-11-21 08:56:11,515 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 21.44 sec
2024-11-21 08:56:17,658 Stage-2 map = 100%, reduce = 75%, Cumulative CPU 33.47 sec
2024-11-21 08:56:20,729 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 37.21 se
c
MapReduce Total cumulative CPU time: 37 seconds 210 msec
Ended Job = job_1732089968849_2334
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 4 Cumulative CPU: 29.25 sec HDFS Read: 3150189 H
DFS Write: 2335868 SUCCESS
Stage-Stage-2: Map: 3 Reduce: 4 Cumulative CPU: 37.21 sec HDFS Read: 3115162 H
DFS Write: 1909 SUCCESS
Total MapReduce CPU Time Spent: 1 minutes 6 seconds 460 msec
OK
Iqaluit Clyde River
Pond Inlet Clyde River
Iqaluit Clyde River
Pond Inlet Clyde River
Wabush Schefferville
Wabush Schefferville
Quebec Jean Lesage Intl Schefferville
Sept Iles Schefferville
Kuujuaq Schefferville
Calgary Intl Red Deer Regional
Time taken: 58.257 seconds, Fetched: 10 row(s)
hive (cdac_roy)>
```

2.)

SPARK-----
Q1.)
1.)

Home

Subscription Details | Nuvepro

Hue - File Browser

cdacuser82315@ip-172-31-9-1

240840325028_Big_data_exam

cdacnppc.cloudloka.com:8132/hue/filebrowser/view=%2Fuser%2Fcdacuser82315?type=hive#user/cdacuser82315/training

GmailMapsLearn Java: Object...Prepare Python Inte...This is my technical...Google IT Automati...GitHub - Python-W...Python Training, Pra...Commonly asked D...All Bookmarks

Search data and saved documents...

Jobs

1sept2024_anup

Tables (5) +

Filter...

airbnb_nyc_data

avro_hive_table

post_office_district

post_office_india

users

File Browser

Search for file name

ActionsDelete forever

UploadNew

Home / user / cdacuser82315 / training

	Name	Size	User	Group	Permissions	Date
<input type="checkbox"/>	.		cdacuser82315	cdacuser82315	drwxr-x--	November 17, 2024 10:14 PM
<input type="checkbox"/>	.		cdacuser82315	cdacuser82315	drwxr-xr-x	November 19, 2024 11:10 AM
<input type="checkbox"/>	100-0.txt	5.5 MB	cdacuser82315	cdacuser82315	-rw-r--	November 12, 2024 05:22 AM
<input type="checkbox"/>	Airlines.csv	271.9 KB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Airports.csv	736.6 KB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Employees.txt	101 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 08:54 PM
<input type="checkbox"/>	NYSE.csv	39.1 MB	cdacuser82315	cdacuser82315	-rw-r--	October 26, 2024 04:34 AM
<input type="checkbox"/>	Routes.csv	2.3 MB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Speed-data.txt	167 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 03, 2024 08:17 PM
<input checked="" type="checkbox"/>	airlines.csv	1.8 KB	cdacuser82315	cdacuser82315	-rw-r--	November 19, 2024 11:10 AM
<input type="checkbox"/>	custs.txt	382.7 KB	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 10:52 PM
<input type="checkbox"/>	desig.txt	54 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 08:54 PM
<input type="checkbox"/>	eclipse.gz	197.5 MB	cdacuser82315	cdacuser82315	-rw-r--	October 26, 2024 04:35 AM

31°C Smoke

Search

ENG IN 2:54 PM 11/21/2024

Home

Subscription Details | Nuvepro

Hue - File Browser

cdacuser82315@ip-172-31-9-1

240840325028_Big_data_exam

cdacnppc.cloudloka.com:8132/hue/filebrowser/view=%2Fuser%2Fcdacuser82315?type=hive#user/cdacuser82315/training

GmailMapsLearn Java: Object...Prepare Python Inte...This is my technical...Google IT Automati...GitHub - Python-W...Python Training, Pra...Commonly asked D...All Bookmarks

Search data and saved documents...

Jobs

1sept2024_anup

Tables (5) +

Filter...

airbnb_nyc_data

avro_hive_table

post_office_district

post_office_india

users

File Browser

Search for file name

ActionsDelete forever

UploadNew

Home / /user/cdacuser82315/training

	Name	Size	User	Group	Permissions	Date
<input type="checkbox"/>	.		cdacuser82315	cdacuser82315	drwxr-x--	November 17, 2024 10:14 PM
<input type="checkbox"/>	.		cdacuser82315	cdacuser82315	drwxr-xr-x	November 19, 2024 11:10 AM
<input type="checkbox"/>	100-0.txt	5.5 MB	cdacuser82315	cdacuser82315	-rw-r--	November 12, 2024 05:22 AM
<input type="checkbox"/>	Airlines.csv	271.9 KB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Airports.csv	736.6 KB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Employees.txt	101 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 08:54 PM
<input type="checkbox"/>	NYSE.csv	39.1 MB	cdacuser82315	cdacuser82315	-rw-r--	October 26, 2024 04:34 AM
<input type="checkbox"/>	Routes.csv	2.3 MB	cdacuser82315	cdacuser82315	-rw-r--	November 06, 2024 10:37 AM
<input type="checkbox"/>	Speed-data.txt	167 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 03, 2024 08:17 PM
<input checked="" type="checkbox"/>	airlines.csv	1.8 KB	cdacuser82315	cdacuser82315	-rw-r--	November 19, 2024 11:10 AM
<input type="checkbox"/>	custs.txt	382.7 KB	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 10:52 PM
<input type="checkbox"/>	desig.txt	54 bytes	cdacuser82315	cdacuser82315	-rw-r--	November 04, 2024 08:54 PM
<input type="checkbox"/>	eclipse.gz	197.5 MB	cdacuser82315	cdacuser82315	-rw-r--	October 26, 2024 04:35 AM

31°C Smoke

Search

ENG IN 2:54 PM 11/21/2024


```
at org.apache.spark.InterruptibleIterator.to(InterruptibleIterator.scala:28)
at scala.collection.TraversableOnce.toBuffer(TraversableOnce.scala:307)
at scala.collection.TraversableOnce.toBuffer$(TraversableOnce.scala:307)
at org.apache.spark.InterruptibleIterator.toBuffer(InterruptibleIterator.scala:28)
at scala.collection.TraversableOnce.toArray(TraversableOnce.scala:294)
at scala.collection.TraversableOnce.toArray$(TraversableOnce.scala:288)
at org.apache.spark.InterruptibleIterator.toArray(InterruptibleIterator.scala:28)
at org.apache.spark.api.python.PythonRDD$.anonfun$runJob$1(PythonRDD.scala:166)
at org.apache.spark.SparkContext$.anonfun$runJob$5(SparkContext.scala:2236)
at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:90)
at org.apache.spark.scheduler.Task.run(Task.scala:131)
at org.apache.spark.executor.Executor$TaskRunner$.anonfun$run$3(Executor.scala:497)
at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:1439)
at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:500)
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
... 1 more

>>> split=eliminate.map(lambda a: (a.split(",")[0],int(a.split(",")[1]),float(a.split(",")[2]),float(a.split(",")[3])))
>>> for line in split.take(10):
...     print(line)
...
('1995', 1, 296.9, 46561.0)
('1995', 2, 296.8, 37443.0)
('1995', 3, 287.51, 34128.0)
('1995', 4, 287.78, 30388.0)
('1996', 1, 283.97, 47808.0)
('1996', 2, 275.78, 43020.0)
('1996', 3, 269.49, 38952.0)
('1996', 4, 278.33, 37443.0)
('1997', 1, 283.4, 35067.0)
('1997', 2, 289.44, 46565.0)
>>>
```

2.)

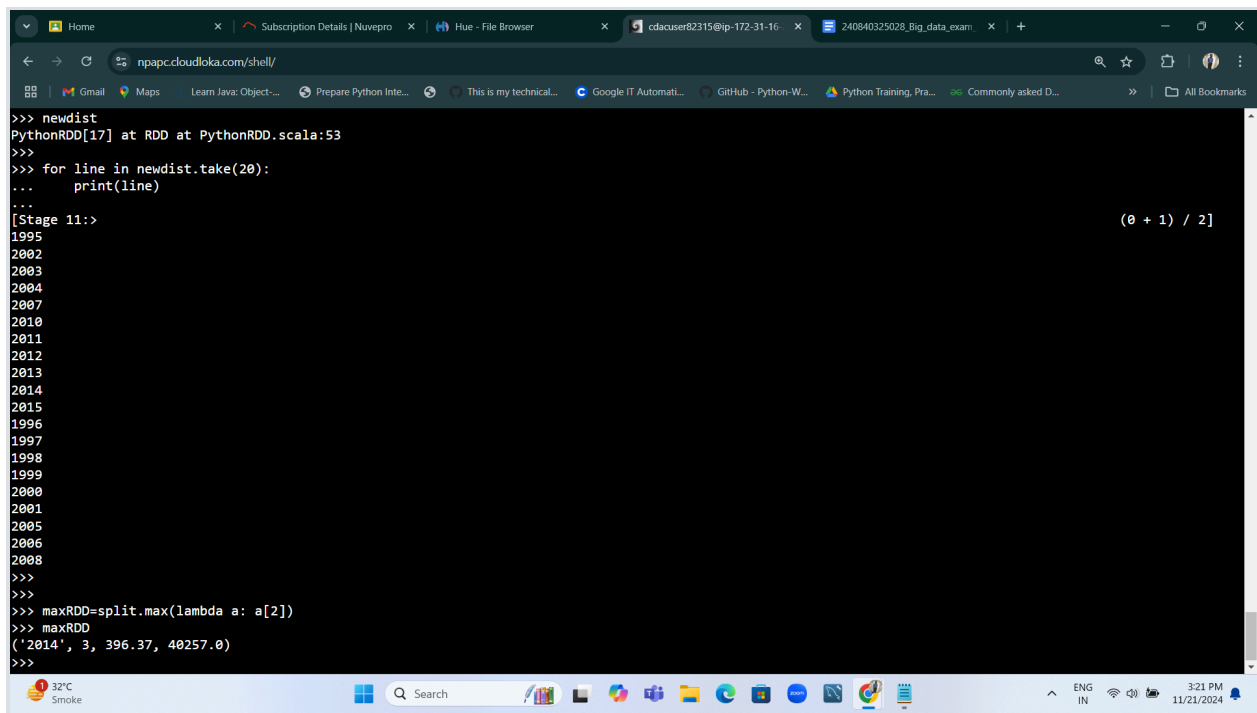
```
False
False
True
>>> newdist=split.map(lambda a: a[0]).distinct()
>>> newdist
PythonRDD[17] at RDD at PythonRDD.scala:53
>>>
>>> for line in newdist.take(20):
...     print(line)
...
[Stage 11:] (0 + 1) / 2]
1995
2002
2003
2004
2007
2010
2011
2012
2013
2014
2015
1996
1997
1998
1999
2000
2001
2005
2006
2008
>>>
>>>
```

Question 2.)

1.)

```
maxRDD=split.max(lambda a: a[2])
```

```
maxRDD
```



The screenshot shows a Jupyter Notebook interface with a dark theme. The browser tabs at the top include 'Home', 'Subscription Details | Nuvepro', 'Hue - File Browser', and 'cdacuser82315@ip-172-31-16'. The address bar shows 'npapc.cloudloka.com/shell/'. The notebook content displays the following code and output:

```
>>> newdist
PythonRDD[17] at RDD at PythonRDD.scala:53
>>>
>>> for line in newdist.take(20):
...     print(line)
...
[Stage 11:>                                     (0 + 1) / 2]
1995
2002
2003
2004
2007
2010
2011
2012
2013
2014
2015
1996
1997
1998
1999
2000
2001
2005
2006
2008
>>>
>>>
>>> maxRDD=split.max(lambda a: a[2])
>>> maxRDD
('2014', 3, 396.37, 40257.0)
>>>
```

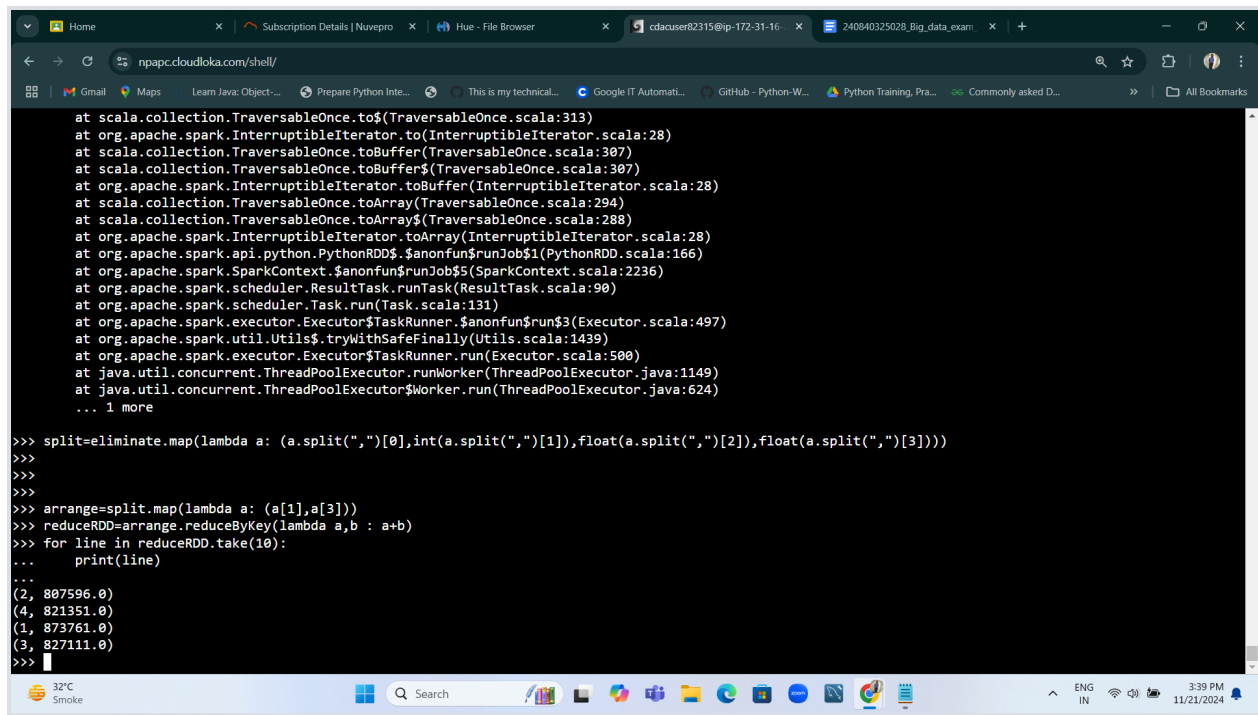
The bottom of the screen shows a Windows taskbar with a search bar, various application icons, and system tray information including '32°C Smoke', 'ENG IN', and the date '11/21/2024'.

```
minRDD=split.min(lambda a: a[2])
```

```
minRDD
```


3.)

```
arrange=split.map(lambda a: (a[1],a[3]))
>>> reduceRDD=arrange.reduceByKey(lambda a,b : a+b)
>>> for line in reduceRDD.take(10):
...     print(line)
```



The screenshot shows a terminal window with a dark background. The top part displays a series of Spark execution logs, including messages from `scala.collection.TraversableOnce`, `org.apache.spark.InterruptibleIterator`, `scala.collection.TraversableOnce`, `org.apache.spark.InterruptibleIterator`, `scala.collection.TraversableOnce`, `org.apache.spark.InterruptibleIterator`, `scala.collection.TraversableOnce`, `org.apache.spark.InterruptibleIterator`, `org.apache.spark.api.python.PythonRDD`, `org.apache.spark.SparkContext`, `org.apache.spark.scheduler.ResultTask`, `org.apache.spark.scheduler.Task`, `org.apache.spark.executor.Executor$TaskRunner`, `org.apache.spark.util.Utils`, `org.apache.spark.executor.Executor$TaskRunner`, `java.util.concurrent.ThreadPoolExecutor`, and `java.util.concurrent.ThreadPoolExecutor`. The logs end with `... 1 more`.

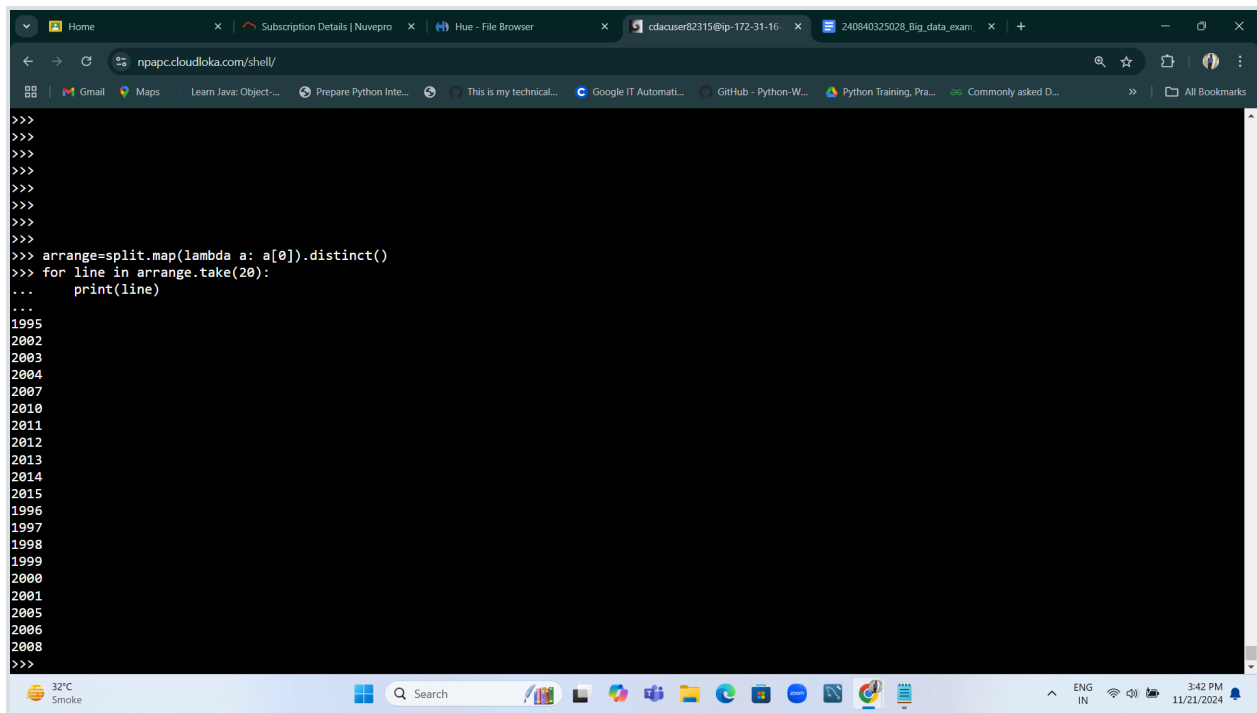
The bottom part of the terminal shows a Jupyter Notebook session with the following code:

```
>>> split=eliminate.map(lambda a: (a.split(",")[0],int(a.split(",")[1]),float(a.split(",")[2]),float(a.split(",")[3])))
>>>
>>>
>>> arrange=split.map(lambda a: (a[1],a[3]))
>>> reduceRDD=arrange.reduceByKey(lambda a,b : a+b)
>>> for line in reduceRDD.take(10):
...     print(line)
...
(2, 807596.0)
(4, 821351.0)
(1, 873761.0)
(3, 827111.0)
>>>
```

The terminal window has a taskbar at the bottom showing the system clock as 3:39 PM on 11/21/2024, and the temperature as 32°C with a 'Smoke' alert.

4.)

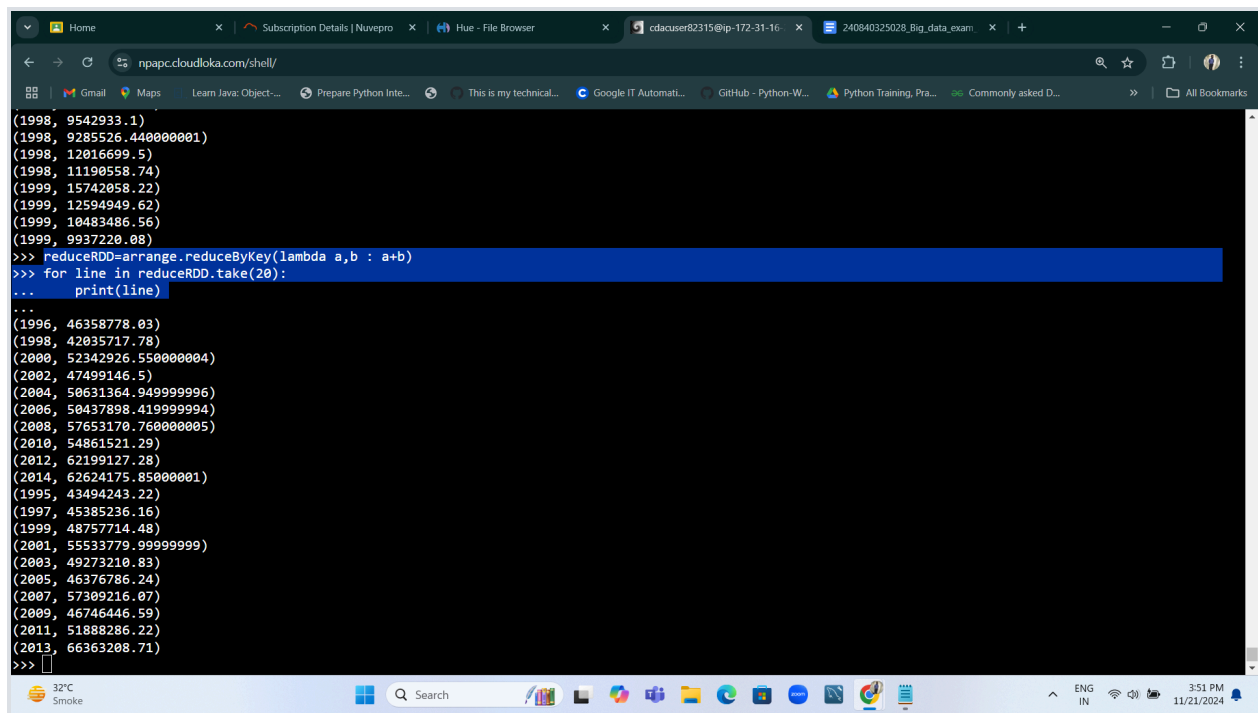
```
arrange=split.map(lambda a: a[0]).distinct()  
>>> for line in arrange.take(20):  
...     print(line)
```



The screenshot shows a web browser window with the address bar displaying `npapc.cloudloka.com/shell/`. The browser's address bar and tabs are visible at the top. The main content area is a terminal window with a black background and white text. The terminal shows the execution of the code from the previous block. The output consists of 20 years listed vertically: 1995, 2002, 2003, 2004, 2007, 2010, 2011, 2012, 2013, 2014, 2015, 1996, 1997, 1998, 1999, 2000, 2001, 2005, 2006, and 2008. The terminal window is part of a desktop environment, with a taskbar at the bottom showing various application icons and system status information like temperature (32°C) and time (3:42 PM, 11/21/2024).

```
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>>  
>>> arrange=split.map(lambda a: a[0]).distinct()  
>>> for line in arrange.take(20):  
...     print(line)  
...  
1995  
2002  
2003  
2004  
2007  
2010  
2011  
2012  
2013  
2014  
2015  
1996  
1997  
1998  
1999  
2000  
2001  
2005  
2006  
2008  
>>>
```

```
5.) reduceRDD=arrange.reduceByKey(lambda a,b : a+b)
>>> for line in reduceRDD.take(20):
...     print(line)
```



The screenshot shows a web browser window with a terminal interface. The terminal displays the execution of a Python script. The code defines a function to calculate the sum of two values for each key in a dataset. The output shows 20 lines of results, each containing a year and a sum. The code is as follows:

```
>>> reduceRDD=arrange.reduceByKey(lambda a,b : a+b)
>>> for line in reduceRDD.take(20):
...     print(line)
...
(1998, 9542933.1)
(1998, 9285526.440000001)
(1998, 12016699.5)
(1998, 11190558.74)
(1999, 15742058.22)
(1999, 12594949.62)
(1999, 10483486.56)
(1999, 9937220.08)
(1996, 46358778.03)
(1998, 42035717.78)
(2000, 52342926.550000004)
(2002, 47499146.5)
(2004, 50631364.949999996)
(2006, 50437898.419999994)
(2008, 57653170.760000005)
(2010, 54861521.29)
(2012, 62199127.28)
(2014, 62624175.85000001)
(1995, 43494243.22)
(1997, 45385236.16)
(1999, 48757714.48)
(2001, 55533779.99999999)
(2003, 49273210.83)
(2005, 46376786.24)
(2007, 57309216.07)
(2009, 46746446.59)
(2011, 51888286.22)
(2013, 66363208.71)
>>>
```

2.) `groupby('year','quarter').agg(count(col(avg_per_seat)>$290)`

→Screen stuck written the query

```
npapc.cloudloka.com/shell/
(1999, 9937720.00)
>>> reduceRDD=arrange.reduceByKey(lambda a,b : a+b)
>>> for line in reduceRDD.take(20):
...     print(line)
...
(1996, 46358778.03)
(1998, 42035717.78)
(2000, 52342926.550000004)
(2002, 47499146.5)
(2004, 50631364.949999996)
(2006, 50437898.419999994)
(2008, 57653170.760000005)
(2010, 54861521.29)
(2012, 62199127.28)
(2014, 62624175.85000001)
(1995, 43494243.22)
(1997, 45385236.16)
(1999, 48757714.48)
(2001, 55533779.99999999)
(2003, 49273210.83)
(2005, 46376786.24)
(2007, 57309216.07)
(2009, 46746446.59)
(2011, 51888286.22)
(2013, 66363208.71)
>>>
>>>
>>>
>>>
>>>
>>>
>>> a
```